

# Persistence

CE881: Mobile and Social Application Programming

Simon Lucas & Spyros Samothrakis

February 23, 2015

- 1 Interesting Cultural Artefacts
- 2 Overall
- 3 Filesystem
- 4 SQL storage
- 5 Discussion

## Theme: “Persistence and Memory”

- Johny Mnemonic (Movie - 1995)
- 320GB of storage was a big thing
- Total Recall

# Main Approaches

- Saving Key-Value pairs in preferences
  - Easy but limited
- Using the File System
  - Android sits on top of a Linux file system
  - There are restrictions on where files can be opened (more of this later)
  - But once you have a `FileInputStream` or a `FileOutputStream`, it is simply standard Java
- Using an SQLite database
  - Lightweight standalone SQL database
  - Standard on Android Platform

## Binary versus Character Data

- jpeg, mp3, Java serialized objects are all binary format
  - This means that each byte in the body of the file can contain any bytes
- Any one with the word Reader or Writer is limited to character data
  - Writing certain bytes will not work: they will be escaped and cause an error in the binary file format
  - Example: try writing a jpeg image using a FileWriter object;
    - It won't work!

# Character Data

- Plain text files
- HTML, XML files
- But: sometimes we need to encode binary data over Character channels
- **Question:** when does this need arise?
  - Solution: use base-64
  - This encodes arbitrary byte sequences using bytes allowed character sequences
  - The cost is that it requires  $1/3$  more space
  - Each 3 bytes of binary require 4 bytes of base 64

## Key-Value Sets

- Like Map structures in Java
- But limited in the range of values they can store
- Cannot store general Objects
- Writing:

```
SharedPreferences sharedPref =  
    getActivity().getPreferences(Context.MODE_PRIVATE);  
SharedPreferences.Editor editor = sharedPref.edit();  
editor.putInt(getString(R.string.saved_high_score), newHighScore);  
editor.commit();
```

## Reading from a Key-Value Set

```
sharedPreferences sharedPreferences =  
    getActivity().getPreferences(Context.MODE_PRIVATE);  
int defaultValue =  
    getResources().getInteger(R.string.saved_high_score_default);  
long highScore =  
    sharedPreferences.getInt(getString(R.string.saved_high_score),  
        defaultValue);
```

Note: can only save simple types and Strings



## Plain text files

- HTML, XML files
- But: sometimes we need to encode binary data over Character channels
- **Question:** when does this need arise?
  - Solution: use base-64
  - This encodes arbitrary byte sequences using bytes allowed character sequences
  - The cost is that it requires  $1/3$  more space
  - Each 3 bytes of binary require 4 bytes of base 64

# File Storage

- First we'll look at some Android specific features of file storage
- This relates to where files can be opened
- And what permissions are required
- Then we'll move on to more general points about file storage
  - In particular, storing data in files with minimal programming effort

# Android File System

- <http://developer.android.com/training/basics/data-storage/files.html>
- Internal versus External Storage
- Internal:
  - By default readable and writable only by this App
  - Other more liberal modes have been deprecated
  - Every app can access its own internal storage
  - No need to request permission in the manifest file
  - All files deleted when an app is removed from a device

## Internal Files No.1 – Direct File Access

### File Creation and Appending:

```
FileOutputStream fos =  
    openFileOutput("test.txt", Context.MODE_PRIVATE); // create new  
FileOutputStream fos =  
    openFileOutput("test.txt", Context.MODE_APPEND); // append
```

- **openFileOutput()** is a method of the Context class
- Activity is a subclass of Context
- Opening for reading:
  - **FileInputStream fis = openFileInput("test.txt");**

## Internal Files No.2

- Alternatively, can make calls to get a File object and then open it for reading, writing or appending in standard Java ways:
  - `File file = new File(context.getFilesDir(), filename);`
- Replace the call to `context.getFilesDir()` with `context.getCacheDir()` for temporary files

## External Storage

- Memory outside of an App's own area
  - May even be on an SD-card
  - Therefore an App cannot guarantee access to it
- Should handle this gracefully!
- Potentially readable/writable by the user and by other apps
- An obvious choice for sharing data

## Using External Storage

- Add the permission to the manifest file:

```
<manifest ...>  
  <uses-permission  
    android:name=  
      "android.permission.WRITE_EXTERNAL_STORAGE" />  
... </manifest>
```

- Replace with .READ\_EXTERNAL\_STORAGE for read-only
- Then use file system in standard Java ways

## More on External Storage

- Check availability before using:

```
/* Checks if external storage  
is available for read and write */  
public boolean isExternalStorageWritable() {  
    String state =  
        Environment.getExternalStorageState();  
    if  
        (Environment.MEDIA_MOUNTED.equals(state)) {  
        return true;  
    }  
    return false;  
}
```

- Similar method for Read Only test – see docs



## Ad Hoc File Formats

- Data is written out in an entirely application specific way
- Conventions are adopted or invented on the fly by the programming team
  - Very flexible
  - Can choose exactly what data to write and how to format it
- Hard work:
  - May need lots of lines of code
  - And careful effort is needed to keep reader and writer in perfect harmony!
- Okay for simple cases, not for complex Apps

# Object Serializers

- Unlike application specific formats, these read and write a wide variety of Object structures
  - In a domain-independent way
  - Some of them may have readers and writers in a variety of languages: hence can exchange object data between different languages
  - If they do what you want, they:
    - Are easy to use
- Involve minimal programming effort
- We'll look at three examples...

# Java's Native Object Serialisation

- If a class that “implements” Serializable
  - Then Objects of that class can automatically be written to and read from Object streams
  - Very easy
  - Fast binary format, low storage space
  - Handles circular references
- But:
  - Can be hard to recover objects if the classes change
  - Restricted to Java
  - Not human readable
  - Binary format cannot directly be sent over text channels

## XML Serializers e.g. WOX

- WOX = Web Objects in XML
- Java version by Lucas (2004) and extended by Jaimez: addition of base-64 for byte arrays and C# readers and writers
- Handles objects of most classes
- XML-based, so not as compact as plain text or as binary
- But given that it's XML, is efficient
- Important: handles circular references

# JSON = JavaScript Object Notation

- Simple lightweight text-based standard for reading-and writing objects
- Efficient and compact
- Supported by MANY languages and platforms
- In many cases the best option except for:
- Binary formats (image, audio, video)
- Object graphs with circular references
- Infuriatingly, JSON cannot handle these
- However, might be worth trying to work around this. . .

## JSON Continued

- I recommend the GSON library from Google for using JSON
- Very easy to use
- However, due to limitations of GSON format there are some cases it does not handle easily
  - E.g. When the declared type of a field is an interface type
  - ( Can customise it to cope with this, but this is extra work
- (WOX handles those cases easily)

# GSON Sample (part 1; from doc)

```
public class GsonTest {  
    public static void main(String[] args) {  
  
        // (Serialization)  
        BagOfPrimitives obj = new BagOfPrimitives();  
        Gson gson = new Gson();  
        String json = gson.toJson(obj);  
        // ==> json is {"value1":1,"value2":"abc"}  
  
        // Note that you can not serialize objects with  
        // circular references since that will result in infinite recursion.  
  
        System.out.println(json);  
        // (Deserialization)  
        BagOfPrimitives obj2 =  
            gson.fromJson(json, BagOfPrimitives.class);  
    }  
}
```

## GSON Test Class: Bag of Primitives (but works for reference types also)

```
static class BagOfPrimitives {  
    private int value1 = 1;  
    private String value2 = "abc";  
    private transient int value3 = 3;  
  
    public static String test = "BOO";  
  
    BagOfPrimitives() {  
        // no-args constructor  
    }  
}
```



# SQL DBs and Android


- If your app needs to store and retrieve data in a flexible way
  - Consider using a relational DB
  - Helps ensure data integrity
  - Well designed relations cannot store self-inconsistent data
- Standard SQL language
  - Very powerful for sorting and selecting the data needed
  - For simple apps using SQL is harder work than simply writing data in JSON format to file system
  - But worth the effort when appropriate

# Relational Modelling

- Relation == table
- Each column specifies the column name and type
- Database Schema (loosely speaking)
- The column names and types for each table
- Each row is a single record of data with a value for each column
- Depending on settings, cell entries may be NULL
- Dilemmas sometimes arise regarding how far to normalise a table

# Normalisation

- How would you model the following Contact DB?
- Each person is identified by name and primary email address
  - Each person may have a number of telephones (home, office, mobile etc.)
  - When designing an App be prepared to compromise:
  - Perfect normalisation versus realistic usage
  - Higher normal forms can sometimes be less efficient



# SQLite

- Separate open-source project
  - On Android platform as standard
  - Hence default choice of relational DB for Android
  - Other choices of pure Java DB also possible
- But would add to size of App
  - Standalone DB
  - Does not run as a separate server process
  - Supports most but not all SQL statements
- Transactions

# Transactions

- Help ensure that DB is always kept in a consistent state
- Each transaction maps the state of the data from one consistent state to the next
- ACID properties

# Atomicity

- A transaction may involve a set or sequence of updates
- Atomicity ensures that either that ALL happen, or NONE of them happen
- Enclose sequence between begin transaction and end transaction statements
- Example: imagine the phone battery died during a transaction
- All the statements executed so far are held in temporary storage
- And only committed at the end (e.g. by moving a file pointer)

# Consistency

- The database is always kept consistent
- Helped by:
  - Suitably high normal form
- Other transactional properties: Atomicity and Isolation

# Isolated

- When concurrent threads are hitting the DB
- Or data is being processed in App using concurrent threads
- Must ensure that threads do not interfere
  - (see example in Threads and Surface Views lecture notes)
- Isolated / Isolation property guarantees this
- May be achieved using locking mechanisms such as Semaphores, or Java statements such as synchronized



# Durable

- An obvious property!
- Once made a transaction should be saved permanently
- And not be affected by systems crashes or power loss

## Package: `android.database.sqlite`

- Contains the SQLite database management classes that an application can use to manage its own private database
- Also has the `sqlite3` database tool in the `tools/` folder
- Use this tool to browse or run SQL commands on the device. Run by typing `sqlite3` in a shell window.

# SQLite versus JDBC


- Android has a set of non-JDBC classes to work directly with SQLite
  - However, JDBC drivers are also available for SQLite
  - See: <https://github.com/SQLDroid/SQLDroid>
  - [http://en.wikibooks.org/wiki/Java\\_JDBC\\_using\\_SQLite/Introduction](http://en.wikibooks.org/wiki/Java_JDBC_using_SQLite/Introduction)
- Hence another possible option would be to use a JDBC driver
- This offers better standardisation, and could be worth exploring
- For these notes we're going to stick with the Android API

## Accessing DBs from Code

- A number of approaches possible
  - Can embed SQL strings in App
  - And make DB calls from wherever needed
- Or:
  - Define a special data handling class
  - All DB access is via handler
  - Main App code only sees objects, never SQL strings
- Or:
  - Can use an automated / semi-automated tool such as Spring / Hibernate
- Discussion question: which way is best?

## DB Interface / Helper Approach

- Define all DB access methods in an interface
- Then provide one or more implementations of this as required e.g.:
- SQLite Implementation: the real thing
  - DummyDB: implement all the methods but don't actually save the data persistently
  - DummyDB can be very useful for testing the App



## Exercise

- Write a DB interface class called ScoreDBInterface to support the storage and retrieval of all scores for a game
- Interface should allow storage of each score as it occurs
- And retrieval of high score, and list of top N scores
- Each score record consists of:
  - Time Stamp of type long
  - Player Name of type String
  - Score achieved of type int
- Write the method signatures and also any convenience classes that you need

## Android Approach

- Similar to DBInterface approach
- Defines data contract class
  - This includes an inner class for each table defining its column names
  - Methods to create tables and drop them
  - Methods to execute queries
- See Android Developer example
  - <http://developer.android.com/training/basics/data-storage/databases.html>
  - Also CE881 Lab code, outlined below

# SQLite in Android

- Construct SQL strings
- Execute SQL strings to perform actions
  - Also use specific helper methods such as query()
- These take arguments to build a SELECT clause with
- Or use.rawQuery() – pass the SQL string
- Both query() and.rawQuery() return a Cursor object
- This is used to iterate over the resulting table of rows



## Score DB Example

- Adapted from FeedReader example:
- <http://developer.android.com/training/basics/data-storage/databases.html>
- ScoreDBContract
  - Class used to define table names
  - Has inner classes for each table
  - In this case, just one table called entry (for entries in the score table)
  - All columns defined in abstract class **ScoreEntry**
- ScoreHelper class
  - Manages access to the DB
  - Declares types for each table column
  - Methods for creating and dropping tables
  - May also implement **ScoreDBInterface**

## Snippets from ScoreHelper

```
private static final String SQL_CREATE_ENTRIES =
    "CREATE TABLE " + TABLE_NAME + " (" +
        ScoreDBContract.ScoreEntry._ID +
        " INTEGER PRIMARY KEY," +
        ScoreDBContract.ScoreEntry.COLUMN_NAME_ENTRY_ID +
        TEXT_TYPE + COMMA_SEP +
        ScoreDBContract.ScoreEntry.COLUMN_NAME_PERSON +
        TEXT_TYPE + COMMA_SEP +
        ScoreDBContract.ScoreEntry.COLUMN_NAME_SCORE +
        "INTEGER" + " ) ";

private static final String SQL_DELETE_ENTRIES =
    "DROP TABLE IF EXISTS " + TABLE_NAME;

public void onCreate(SQLiteDatabase db) {
    db.execSQL(SQL_CREATE_ENTRIES);
}
```

## Adding an Entry...

```
SQLiteDatabase db =  
    scoreHelper.getWritableDatabase();  
scoreHelper.addEntry(db, "Simon", "" +  
    10 * (1 + random.nextInt(100)));  
Log.i(TAG, "Populated Table");  
db.close();
```

## addEntry Method

```
public void addEntry(SQLiteDatabase db, String person, String score) {  
    // Create a new map of values, where column names are the keys  
    ContentValues values = new ContentValues();  
    values.put(ScoreDBContract.ScoreEntry.COLUMN_NAME_ENTRY_ID, id++);  
    values.put(ScoreDBContract.ScoreEntry.COLUMN_NAME_PERSON, person);  
    values.put(ScoreDBContract.ScoreEntry.COLUMN_NAME_SCORE, score);  
  
    // Insert the new row, returning the primary key value of the new row  
    long newRowId;  
    newRowId = db.insert(  
        ScoreDBContract.ScoreEntry.TABLE_NAME,  
        ScoreDBContract.ScoreEntry.COLUMN_NAME_NULLABLE,  
        values);  
    Log.i(SQLiteActivity.TAG, "Inserted row: " + newRowId);  
}
```

## Using a Cursor with a Query Selecting all scores

```
Cursor cursor = db.rawQuery("Select * from " +  
    ScoreDBContract.ScoreEntry.TABLE_NAME,  
    new String[] {});  
cursor.moveToFirst();  
// need to find index of each column before retrieving it  
int scoreIndex =  
    cursor.getColumnIndex(ScoreDBContract.ScoreEntry.COLUMN_NAME_SCORE);  
while (!cursor.isAfterLast()) {  
    int score = cursor.getInt(scoreIndex);  
    boolean flag = cursor.moveToNext();  
    Log.i(TAG, flag + " : " + score);  
}
```

## More details...

- See SQLite lab exercise
  - Including .zip with all the details
  - The example opens and closes a DB Connection each time it is needed
  - This is perhaps not the most efficient way
  - But it saves thinking through lifecycle methods
- More efficient way:
  - open connection when activity is created
  - Close connection when activity is destroyed

## Summary (1)

- Data serialization is an important topic
- These notes discussed storage and retrieval on file systems
- But much of this same applies for sending over a network
- Choose carefully: when given a free choice JSON is a good default

## Summary (2)

- Android ships with SQLite database
- Natural choice for storing data that must be flexibly queried
- And have guaranteed integrity
- However, API is rather low level
- Best approach is to embed all DB access code in a separate helper class (or classes)
- My favourite: all access goes through a DB interface class
- Can then test rest of app using a dummy implementation
- Work through (next week's) lab exercise