

Persistence

CE881: Mobile and Social Application Programming

Spyros Samothrakis

February 22, 2016

Interesting Cultural Artefacts

Persistence

Filesystem

SQL storage

Discussion

THEME: “PERSISTENCE AND MEMORY”

- ▶ Johny Mnemonic (Movie - 1995)
- ▶ 320GB of storage was a big thing
- ▶ Total Recall

PRESENTATIONS

- ▶ Week 8 (i.e. next week)
 - ▶ No need to prepare anything
 - ▶ Short, 5 minutes presentation on what you do
- ▶ Week 10 (lab 10) short presentations/ questions
- ▶ Assiment 2 to be released shortly

MAIN APPROACHES

- ▶ Saving Key-Value pairs in preferences
 - ▶ Easy but limited
- ▶ Using the File System
 - ▶ Android sits on top of a Linux file system
 - ▶ There are restrictions on where files can be opened (more of this later)
 - ▶ But once you have a `FileInputStream` or a `FileOutputStream`, it is simply standard Java
- ▶ Using an SQLite database
 - ▶ Lightweight standalone SQL database
 - ▶ Standard on Android Platform

BINARY VERSUS CHARACTER DATA

- ▶ jpeg, mp3, Java serialized objects are all binary format
 - ▶ This means that each byte in the body of the file can contain any bytes
- ▶ Any one with the word Reader or Writer is limited to character data
 - ▶ Writing certain bytes will not work: they will be escaped and cause an error in the binary file format
 - ▶ Example: try writing a jpeg image using a FileWriter object;
 - ▶ It won't work!

CHARACTER DATA

- ▶ Plain text files
- ▶ HTML, XML files
- ▶ But: sometimes we need to encode binary data over Character channels
- ▶ **Question:** when does this need arise?
 - ▶ Solution: use base-64
 - ▶ This encodes arbitrary byte sequences using bytes allowed character sequences
 - ▶ The cost is that it requires 1/3 more space
 - ▶ Each 3 bytes of binary require 4 bytes of base 64

KEY-VALUE SETS

- ▶ Like Map structures in Java
- ▶ But limited in the range of values they can store
- ▶ Cannot store general Objects
- ▶ Writing:

```
SharedPreferences sharedPref =  
    getActivity().getPreferences(Context.MODE_PRIVATE);  
SharedPreferences.Editor editor = sharedPref.edit();  
editor.putInt(getString(R.string.saved_high_score), newHighScore)  
editor.commit()
```


READING FROM A KEY-VALUE SET

```
haredPreferences sharedPref =  
    getActivity().getPreferences(Context.MODE_PRIVATE);  
int defaultValue =  
    getResources().getInteger(R.string.saved_high_score_default);  
long highScore =  
    sharedPref.getInt(getString(R.string.saved_high_score),  
        defaultValue);
```

Note: can only save simple types and Strings

PLAIN TEXT FILES

- ▶ HTML, XML files
- ▶ But: sometimes we need to encode binary data over Character channels
- ▶ **Question:** when does this need arise?
 - ▶ Solution: use base-64
 - ▶ This encodes arbitrary byte sequences using bytes allowed character sequences
 - ▶ The cost is that it requires $1/3$ more space
 - ▶ Each 3 bytes of binary require 4 bytes of base 64

FILE STORAGE

- ▶ First we'll look at some Android specific features of file storage
- ▶ This relates to where files can be opened
- ▶ And what permissions are required
- ▶ Then we'll move on to more general points about file storage
 - ▶ In particular, storing data in files with minimal programming effort

ANDROID FILE SYSTEM

- ▶ <http://developer.android.com/training/basics/data-storage/files.html>
- ▶ Internal versus External Storage
- ▶ Internal:
 - ▶ By default readable and writable only by this App
 - ▶ Other more liberal modes have been deprecated
 - ▶ Every app can access its own internal storage
 - ▶ No need to request permission in the manifest file
 - ▶ All files deleted when an app is removed from a device

INTERNAL FILES No.1 – DIRECT FILE ACCESS

File Creation and Appending:

```
FileOutputStream fos =  
    openFileOutput("test.txt", Context.MODE_PRIVATE); // create new  
FileOutputStream fos =  
    openFileOutput("test.txt", Context.MODE_APPEND); // append
```

- ▶ **openFileOutput()** is a method of the Context class
- ▶ Activity is a subclass of Context
- ▶ Opening for reading:
 - ▶ **FileInputStream fis = openFileInput("test.txt");**

INTERNAL FILES NO.2

- ▶ Alternatively, can make calls to get a File object and then open it for reading, writing or appending in standard Java ways:
 - ▶ `File file = new File(context.getFilesDir(), filename);`
- ▶ Replace the call to `context.getFilesDir()` with `context.getCacheDir()` for temporary files

EXTERNAL STORAGE

- ▶ Memory outside of an App's own area
 - ▶ May even be on an SD-card
 - ▶ Therefore an App cannot guarantee access to it
- ▶ Should handle this gracefully!
- ▶ Potentially readable/writable by the user and by other apps
- ▶ An obvious choice for sharing data

USING EXTERNAL STORAGE

- ▶ Add the permission to the manifest file:

```
<manifest ...>  
  <uses-permission  
    android:name=  
      "android.permission.WRITE_EXTERNAL_STORAGE" />  
... </manifest>
```

- ▶ Replace with `.READ_EXTERNAL_STORAGE` for read-only
- ▶ Then use file system in standard Java ways

MORE ON EXTERNAL STORAGE

- Check availability before using:

```
/* Checks if external storage  
is available for read and write */  
public boolean isExternalStorageWritable() {  
    String state =  
        Environment.getExternalStorageState();  
    if  
        (Environment.MEDIA_MOUNTED.equals(state)) {  
        return true;  
    }  
    return false;  
}
```

- Similar method for Read Only test – see docs

AD HOC FILE FORMATS

- ▶ Data is written out in an entirely application specific way
- ▶ Conventions are adopted or invented on the fly by the programming team
 - ▶ Very flexible
 - ▶ Can choose exactly what data to write and how to format it
- ▶ Hard work:
 - ▶ May need lots of lines of code
 - ▶ And careful effort is needed to keep reader and writer in perfect harmony!
- ▶ Okay for simple cases, not for complex Apps

OBJECT SERIALIZERS

- ▶ Unlike application specific formats, these read and write a wide variety of Object structures
 - ▶ In a domain-independent way
 - ▶ Some of them may have readers and writers in a variety of languages: hence can exchange object data between different languages
 - ▶ If they do what you want, they:
 - ▶ Are easy to use
- ▶ Involves minimal programming effort
- ▶ We'll look at three examples...

JAVA'S NATIVE OBJECT SERIALISATION

- ▶ If a class that “implements” Serializable
 - ▶ Then Objects of that class can automatically be written to and read from Object streams
 - ▶ Very easy
 - ▶ Fast binary format, low storage space
 - ▶ Handles circular references
- ▶ But:
 - ▶ Can be hard to recover objects if the classes change
 - ▶ Restricted to Java
 - ▶ Not human readable
 - ▶ Binary format cannot directly be sent over text channels

XML SERIALIZERS E.G. WOX

- ▶ WOX = Web Objects in XML
- ▶ Java version by Lucas (2004) and extended by Jaimez:
addition of base-64 for byte arrays and C# readers and writers
- ▶ Handles objects of most classes
- ▶ XML-based, so not as compact as plain text or as binary
- ▶ But given that it's XML, is efficient
- ▶ Important: handles circular references

JSON = JAVASCRIPT OBJECT NOTATION

- ▶ Simple lightweight text-based standard for reading-and writing objects
- ▶ Efficient and compact
- ▶ Supported by MANY languages and platforms
- ▶ In many cases the best option except for:
- ▶ Binary formats (image, audio, video)
- ▶ Object graphs with circular references
- ▶ Infuriatingly, JSON cannot handle these
- ▶ However, might be worth trying to work around this...


JSON CONTINUED

- ▶ I recommend the GSON library from Google for using JSON
- ▶ Very easy to use
- ▶ However, due to limitations of GSON format there are some cases it does not handle easily
 - ▶ E.g. When the declared type of a field is an interface type
 - ▶ (Can customise it to cope with this, but this is extra work
- ▶ (WOX handles those cases easily)

GSON SAMPLE (PART 1; FROM DOC)

```
public class GsonTest {  
    public static void main(String[] args) {  
  
        // (Serialization)  
        BagOfPrimitives obj = new BagOfPrimitives();  
        Gson gson = new Gson();  
        String json = gson.toJson(obj);  
        // ==> json is {"value1":1,"value2":"abc"}  
  
        // Note that you can not serialize objects with  
        // circular references since that will result in infinite recursion.  
  
        System.out.println(json);  
        // (Deserialization)  
        BagOfPrimitives obj2 =  
            gson.fromJson(json, BagOfPrimitives.class);  
    }  
}
```


GSON TEST CLASS: BAG OF PRIMITIVES (BUT WORKS FOR REFERENCE TYPES ALSO)



```
static class BagOfPrimitives {  
    private int value1 = 1;  
    private String value2 = "abc";  
    private transient int value3 = 3;  
  
    public static String test = "BOO";  
  
    BagOfPrimitives() {  
        // no-args constructor  
    }  
}
```

SQL DBs AND ANDROID


- ▶ If your app needs to store and retrieve data in a flexible way
 - ▶ Consider using a relational DB
 - ▶ Helps ensure data integrity
 - ▶ Well designed relations cannot store self-inconsistent data
- ▶ Standard SQL language
 - ▶ Very powerful for sorting and selecting the data needed
 - ▶ For simple apps using SQL is harder work than simply writing data in JSON format to file system
 - ▶ But worth the effort when appropriate

RELATIONAL MODELLING

- ▶ Relation == table
- ▶ Each column specifies the column name and type
- ▶ Database Schema (loosely speaking)
- ▶ The column names and types for each table
- ▶ Each row is a single record of data with a value for each column
- ▶ Depending on settings, cell entries may be NULL
- ▶ Dilemmas sometimes arise regarding how far to normalise a table

NORMALISATION

- ▶ How would you model the following Contact DB?
- ▶ Each person is identified by name and primary email address
 - ▶ Each person may have a number of telephones (home, office, mobile etc.)
 - ▶ When designing an App be prepared to compromise:
 - ▶ Perfect normalisation versus realistic usage
 - ▶ Higher normal forms can sometimes be less efficient



SQLite

- ▶ Separate open-source project
 - ▶ On Android platform as standard
 - ▶ Hence default choice of relational DB for Android
 - ▶ Other choices of pure Java DB also possible
- ▶ But would add to size of App
 - ▶ Standalone DB
 - ▶ Does not run as a separate server process
 - ▶ Supports most but not all SQL statements
- ▶ Transactions

TRANSACTIONS

- ▶ Help ensure that DB is always kept in a consistent state
- ▶ Each transaction maps the state of the data from one consistent state to the next
- ▶ ACID properties

ATOMICITY

- ▶ A transaction may involve a set or sequence of updates
- ▶ Atomicity ensures that either that ALL happen, or NONE of them happen
- ▶ Enclose sequence between begin transaction and end transaction statements
- ▶ Example: imagine the phone battery died during a transaction
- ▶ All the statements executed so far are held in temporary storage
- ▶ And only committed at the end (e.g. by moving a file pointer)

CONSISTENCY

- ▶ The database is always kept consistent
- ▶ Helped by:
 - ▶ Suitably high normal form
- ▶ Other transactional properties: Atomicity and Isolation

ISOLATED

- ▶ When concurrent threads are hitting the DB
- ▶ Or data is being processed in App using concurrent threads
- ▶ Must ensure that threads do not interfere
 - ▶ (see example in Threads and Surface Views lecture notes)
- ▶ Isolated / Isolation property guarantees this
- ▶ May be achieved using locking mechanisms such as Semaphores, or Java statements such as synchronized

DURABLE

- ▶ An obvious property!
- ▶ Once made a transaction should be saved permanently
- ▶ And not be affected by systems crashes or power loss

PACKAGE: ANDROID.DATABASE.SQLITE

- ▶ Contains the SQLite database management classes that an application can use to manage its own private database
- ▶ Also has the sqlite3 database tool in the tools/ folder
- ▶ Use this tool to browse or run SQL commands on the device. Run by typing sqlite3 in a shell window.

SQLITE VERSUS JDBC


- ▶ Android has a set of non-JDBC classes to work directly with SQLite
 - ▶ However, JDBC drivers are also available for SQLite
 - ▶ See: <https://github.com/SQLDroid/SQLDroid>
 - ▶ http://en.wikibooks.org/wiki/Java_JDBC_using_SQLite/Introduction
- ▶ Hence another possible option would be to use a JDBC driver
- ▶ This offers better standardisation, and could be worth exploring
- ▶ For these notes we're going to stick with the Android API

ACCESSING DBs FROM CODE

- ▶ A number of approaches possible
 - ▶ Can embed SQL strings in App
 - ▶ And make DB calls from wherever needed
- ▶ Or:
 - ▶ Define a special data handling class
 - ▶ All DB access is via handler
 - ▶ Main App code only sees objects, never SQL strings
- ▶ Or:
 - ▶ Can use an automated / semi-automated tool such as Spring / Hibernate
- ▶ Discussion question: which way is best?

DB INTERFACE / HELPER APPROACH

- ▶ Define all DB access methods in an interface
- ▶ Then provide one or more implementations of this as required e.g.:
- ▶ SQLite Implementation: the real thing
 - ▶ DummyDB: implement all the methods but don't actually save the data persistently
 - ▶ DummyDB can be very useful for testing the App



EXERCISE

- ▶ Write a DB interface class called ScoreDBInterface to support the storage and retrieval of all scores for a game
- ▶ Interface should allow storage of each score as it occurs
- ▶ And retrieval of high score, and list of top N scores
- ▶ Each score record consists of:
 - ▶ Time Stamp of type long
 - ▶ Player Name of type String
 - ▶ Score achieved of type int
- ▶ Write the method signatures and also any convenience classes that you need

ANDROID APPROACH

- ▶ Similar to DBInterface approach
- ▶ Defines data contract class
 - ▶ This includes an inner class for each table defining its column names
 - ▶ Methods to create tables and drop them
 - ▶ Methods to execute queries
- ▶ See Android Developer example
 - ▶ <http://developer.android.com/training/basics/data-storage/databases.html>
 - ▶ Also CE881 Lab code, outlined below

SQLITE IN ANDROID

- ▶ Construct SQL strings
- ▶ Execute SQL strings to perform actions
 - ▶ Also use specific helper methods such as `query()`
- ▶ These take arguments to build a `SELECT` clause with
- ▶ Or use `rawQuery()` – pass the SQL string
- ▶ Both `query()` and `rawQuery()` return a `Cursor` object
- ▶ This is used to iterate over the resulting table of rows

SCORE DB EXAMPLE

- ▶ Adapted from FeedReader example:
- ▶ <http://developer.android.com/training/basics/data-storage/databases.html>
- ▶ ScoreDBContract
 - ▶ Class used to define table names
 - ▶ Has inner classes for each table
 - ▶ In this case, just one table called entry (for entries in the score table)
 - ▶ All columns defined in abstract class **ScoreEntry**
- ▶ ScoreHelper class
 - ▶ Manages access to the DB
 - ▶ Declares types for each table column
 - ▶ Methods for creating and dropping tables
 - ▶ May also implement **ScoreDBInterface**

SNIPPETS FROM SCOREHELPER

```
private static final String SQL_CREATE_ENTRIES =
    "CREATE TABLE " + TABLE_NAME + " (" +
        ScoreDBContract.ScoreEntry._ID +
        " INTEGER PRIMARY KEY," +
        ScoreDBContract.ScoreEntry.COLUMN_NAME_ENTRY_ID +
        TEXT_TYPE + COMMA_SEP +
        ScoreDBContract.ScoreEntry.COLUMN_NAME_PERSON +
        TEXT_TYPE + COMMA_SEP +
        ScoreDBContract.ScoreEntry.COLUMN_NAME_SCORE +
        "INTEGER" + " ) ";

private static final String SQL_DELETE_ENTRIES =
    "DROP TABLE IF EXISTS " + TABLE_NAME;

public void onCreate(SQLiteDatabase db) {
    db.execSQL(SQL_CREATE_ENTRIES);
}
```

ADDING AN ENTRY...

```
SQLiteDatabase db =  
    scoreHelper.getWritableDatabase();  
scoreHelper.addEntry(db, "Simon", "" +  
    10 * (1 + random.nextInt(100)));  
Log.i(TAG, "Populated Table");  
db.close();
```

ADDENTRY METHOD

```
public void addEntry(SQLiteDatabase db, String person, String score) {  
    // Create a new map of values, where column names are the keys  
    ContentValues values = new ContentValues();  
    values.put(ScoreDBContract.ScoreEntry.COLUMN_NAME_ENTRY_ID, id++);  
    values.put(ScoreDBContract.ScoreEntry.COLUMN_NAME_PERSON, person);  
    values.put(ScoreDBContract.ScoreEntry.COLUMN_NAME_SCORE, score);  
  
    // Insert the new row, returning the primary key value of the new row  
    long newRowId;  
    newRowId = db.insert(  
        ScoreDBContract.ScoreEntry.TABLE_NAME,  
        ScoreDBContract.ScoreEntry.COLUMN_NAME_NULLABLE,  
        values);  
    Log.i(SQLiteActivity.TAG, "Inserted row: " + newRowId);  
}
```

USING A CURSOR WITH A QUERY SELECTING ALL SCORES

```
Cursor cursor = db.rawQuery("Select * from " +
    ScoreDBContract.ScoreEntry.TABLE_NAME,
    new String[]{});
cursor.moveToFirst();
// need to find index of each column before retrieving it
int scoreIndex =
    cursor.getColumnIndex(ScoreDBContract.ScoreEntry.COLUMN_NAME_SCORE);
while (!cursor.isAfterLast()) {
    int score = cursor.getInt(scoreIndex);
    boolean flag = cursor.moveToNext();
    Log.i(TAG, flag + " : " + score);
}
```

MORE DETAILS...

- ▶ See SQLite lab exercise
 - ▶ Including .zip with all the details
 - ▶ The example opens and closes a DB Connection each time it is needed
 - ▶ This is perhaps not the most efficient way
 - ▶ But it saves thinking through lifecycle methods
- ▶ More efficient way:
 - ▶ open connection when activity is created
 - ▶ Close connection when activity is destroyed

SUMMARY (1)

- ▶ Data serialization is an important topic
- ▶ These notes discussed storage and retrieval on file systems
- ▶ But much of this same applies for sending over a network
- ▶ Choose carefully: when given a free choice JSON is a good default

SUMMARY (2)

- ▶ Android ships with SQLite database
- ▶ Natural choice for storing data that must be flexibly queried
- ▶ And have guaranteed integrity
- ▶ However, API is rather low level
- ▶ Best approach is to embed all DB access code in a separate helper class (or classes)
- ▶ My favourite: all access goes through a DB interface class
- ▶ Can then test rest of app using a dummy implementation
- ▶ Work through this week's lab exercise
- ▶ Some slides by Simon Lucas